

## Lab 12 : Horizontal Pod Autoscaling

1. Create a Deployment named scaler-challenge.
  - a. Use the image as **registry.k8s.io/hpa-example** and port **80**.
  - b. Define a CPU request of 100m and a limit of 200m.
  - c. Expose it via a Service named scaler-service on port **80**.

workload.yaml

Shell

```
apiVersion: v1
kind: Service
metadata:
  name: scaler-service
spec:
  type: ClusterIP
  selector:
    app: scaler-challenge
  ports:
  - port: 80
    targetPort: 80
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: scaler-challenge
spec:
  selector:
    matchLabels:
      app: scaler-challenge
  template:
    metadata:
      labels:
        app: scaler-challenge
    spec:
      containers:
      - name: scaler-challenge
        image: registry.k8s.io/hpa-example
        resources:
          requests:
            cpu: "100m"
          limits:
            cpu: "200m"
```

```
  ports:
    - containerPort: 80
```

Shell

```
k apply -f workload.yaml
service/scaler-service created
deployment.apps/scaler-challenge created
```

## 2. Create the HPA.

- Metrics: Average CPU Utilization at 60%.

hpa.yaml

```
Shell
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: scaler-challenge-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: scaler-challenge
  minReplicas: 2
  maxReplicas: 6
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 60
```

Shell

```
k apply -f hpa.yaml
horizontalpodautoscaler.autoscaling/scaler-challenge-hpa created
```

Shell

```
k get hpa
NAME           REFERENCE          TARGETS          MINPODS
MAXPODS   REPLICAS   AGE
scaler-challenge-hpa   Deployment/scaler-challenge   cpu: 1%/60%   2           6
2           58s
```

3. Stress the workload.

- Open a terminal and watch the HPA.

Shell

```
kubectl get hpa scaler-challenge-hpa -w
```

- Run the stress test.

Shell

```
kubectl run load-gen --image=busybox:1.28 --restart=Never -- /bin/sh -c "while
true; do wget -q -O- http://scaler-service; done"
```

4. Answer the questions below:

- Why did the Pod count immediately jump to 2 even before you started the load generator?
  - Min Replicas = 2**
- What was the highest number of replicas reached during the stress test?
  - 6 replicas.**
- After you delete the load-gen pod, how long does it take for the replicas to scale back down to 2?
  - Around 5 minutes.**

Shell

```
k get deploy
NAME             READY   UP-TO-DATE   AVAILABLE   AGE
scaler-challenge   6/6       6            6           17m
```

Shell

```
kubectl get hpa scaler-challenge-hpa -w
NAME                  REFERENCE          TARGETS          MINPODS
MAXPODS   REPLICAS   AGE
scaler-challenge-hpa   Deployment/scaler-challenge   cpu: 1%/60%   2        6
2          11m
scaler-challenge-hpa   Deployment/scaler-challenge   cpu: 53%/60%   2        6
2          12m

scaler-challenge-hpa   Deployment/scaler-challenge   cpu: 254%/60%   2
6          2          13m
scaler-challenge-hpa   Deployment/scaler-challenge   cpu: 254%/60%   2
6          4          13m
scaler-challenge-hpa   Deployment/scaler-challenge   cpu: 254%/60%   2
6          6          13m
scaler-challenge-hpa   Deployment/scaler-challenge   cpu: 87%/60%    2
6          6          14m
```

Shell

```
k delete pod load-gen
```